

Advanced Data Analytics

(CSE4029)

Black Friday Sales Analysis

Project Report

(Phase - II)



Varad Vijay Temghare
19BCI7045

Submitted to :
Dr. Nitesh Funde

Introduction:-

In Data Science, we usually deal with datasets that contain multiple labels in one or more than one column. These labels can be in the form of words or numbers. To make the data understandable or in human-readable form, the training data is often labeled in words.

Dataset Link :- [Here](#)

Dimensions of the above dataset:- We have 5,50,069 rows and 12 columns .

Code Link:- [Here](#)

Label Encoding refers to converting the labels into a numeric form so as to convert them into the machine-readable form. Machine learning algorithms can then decide in a better way how those labels must be operated. It is an important preprocessing step for the structured dataset in supervised learning.

```
Encoding the categorical variables

[45] from sklearn.preprocessing import LabelEncoder
     lr = LabelEncoder()

[46] df['Gender'] = lr.fit_transform(df['Gender'])

[47] df['Age'] = lr.fit_transform(df['Age'])

[48] df['City_Category'] = lr.fit_transform(df['City_Category'])

[49] df.head()

  User_ID  Product_ID  Gender  Age  Occupation  City_Category  Marital_Status  Product_Category_1  Product_Category_2  Product_Category_3  Purchase  Stay_In_Current_City_Years_0  Stay_In_Current_City_Years_1
0  1000001  P00069042      0    0      10          0          0          3          NaN          NaN          8370              0
1  1000001  P00248942      0    0      10          0          0          1          6.0          14.0          15200              0
2  1000001  P00087842      0    0      10          0          0          12          NaN          NaN          1422              0
3  1000001  P00085442      0    0      10          0          0          12          14.0          NaN          1057              0
4  1000002  P00285442      1    6      16          2          0          8          NaN          NaN          7969              0

[50] df['Product_Category_2'] = df['Product_Category_2'].fillna(0).astype('int64')
     df['Product_Category_3'] = df['Product_Category_3'].fillna(0).astype('int64')
```

```
[50] df['Product_Category_2'] =df['Product_Category_2'].fillna(0).astype('int64')
df['Product_Category_3'] =df['Product_Category_3'].fillna(0).astype('int64')
```

```
[51] df.isnull().sum()
```

```
User_ID          0
Product_ID       0
Gender           0
Age              0
Occupation       0
City_Category    0
Marital_Status   0
Product_Category_1  0
Product_Category_2  0
Product_Category_3  0
Purchase         0
Stay_In_Current_City_Years_0  0
Stay_In_Current_City_Years_1  0
Stay_In_Current_City_Years_2  0
Stay_In_Current_City_Years_3  0
Stay_In_Current_City_Years_4+  0
dtype: int64
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 16 columns):
#   Column                                     Non-Null Count  Dtype
---  ---
0   User_ID                                   550068 non-null  int64
1   Product_ID                               550068 non-null  object
2   Gender                                   550068 non-null  int64
3   Age                                       550068 non-null  int64
4   Occupation                               550068 non-null  int64
5   City_Category                             550068 non-null  int64
6   Marital_Status                           550068 non-null  int64
7   Product_Category_1                       550068 non-null  int64
8   Product_Category_2                       550068 non-null  int64
9   Product_Category_3                       550068 non-null  int64
10  Purchase                                  550068 non-null  int64
11  Stay_In_Current_City_Years_0             550068 non-null  uint8
12  Stay_In_Current_City_Years_1             550068 non-null  uint8
13  Stay_In_Current_City_Years_2             550068 non-null  uint8
14  Stay_In_Current_City_Years_3             550068 non-null  uint8
15  Stay_In_Current_City_Years_4+            550068 non-null  uint8
dtypes: int64(10), object(1), uint8(5)
memory usage: 48.8+ MB
```

▼ Dropping the irrelevant columns

```
[53] df = df.drop(["User_ID","Product_ID"],axis=1)
```

▼ Splitting data into independent and dependent variables

```
[54] X = df.drop("Purchase",axis=1)
```

```
[55] y=df["Purchase"]
```

```
[56] from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=123)
```

We have split the dataset as ,

TEST DATA -> 30%

TRAIN DATA -> 70%

1. Linear Regression :

```
Modeling

Linear Regression

[57] from sklearn.linear_model import LinearRegression

[58] lr = LinearRegression()
      lr.fit(X_train,y_train)

      LinearRegression()

[59] lr.intercept_

      9536.400764131593

[60] lr.coef_

      array([[ 465.82318446,  112.36643445,   5.05508596,  314.06766138,
        -58.23217776, -348.4514785 ,  12.98415047,  143.49190467,
        -20.83796687,   5.4676518 ,  17.68367185,  -3.96751734,
         1.65416056]])

[61] y_pred = lr.predict(X_test)

[62] from sklearn.metrics import mean_absolute_error,mean_squared_error, r2_score

[63] mean_absolute_error(y_test, y_pred)

      3532.069226165843

[64] mean_squared_error(y_test, y_pred)

      21397853.26940751

[65] r2_score(y_test, y_pred)

      0.15192944521481688

[66] from math import sqrt
      print("RMSE of Linear Regression Model is ",sqrt(mean_squared_error(y_test, y_pred)))

      RMSE of Linear Regression Model is  4625.781368526566
```

2. Decision Tree Regressor :

▼ DecisionTreeRegressor

```
[67] from sklearn.tree import DecisionTreeRegressor
      # create a regressor object
      regressor = DecisionTreeRegressor(random_state = 0)

[68] regressor.fit(X_train, y_train)
      DecisionTreeRegressor(random_state=0)

[69] dt_y_pred = regressor.predict(X_test)

[70] mean_absolute_error(y_test, dt_y_pred)
      2372.0357559134654

[71] mean_squared_error(y_test, dt_y_pred)
      11300579.466797074

[72] r2_score(y_test, dt_y_pred)
      0.5521191505924365

[73] from math import sqrt
      print("RMSE of Linear Regression Model is ",sqrt(mean_squared_error(y_test, dt_y_pred)))
      RMSE of Linear Regression Model is  3361.633452177241
```

3. Random Forest Regressor :

▼ Random Forest Regressor

```
[74] from sklearn.ensemble import RandomForestRegressor
      # create a regressor object
      RFRegressor = RandomForestRegressor(random_state = 0)

[75] RFRegressor.fit(X_train, y_train)
      RandomForestRegressor(random_state=0)

[76] rf_y_pred = RFRegressor.predict(X_test)

[77] mean_absolute_error(y_test, rf_y_pred)
      2222.049109204734

[78] mean_squared_error(y_test, rf_y_pred)
      9310769.87311957

[79] r2_score(y_test, rf_y_pred)
      0.6309821516972987

[80] from math import sqrt
      print("RMSE of Linear Regression Model is ",sqrt(mean_squared_error(y_test, rf_y_pred)))
      RMSE of Linear Regression Model is  3051.35541573242
```

4. XGBoost Regressor :

▼ XGBoost Regressor

```
✓ [81] from xgboost.sklearn import XGBRegressor
```

```
✓ [82] xgb_reg = XGBRegressor(learning_rate=1.0, max_depth=6, min_child_weight=40, seed=0)
xgb_reg.fit(X_train, y_train)
```

[09:40:59] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.
XGBRegressor(learning_rate=1.0, max_depth=6, min_child_weight=40, seed=0)

```
✓ [83] xgb_y_pred = xgb_reg.predict(X_test)
```

```
✓ [84] mean_absolute_error(y_test, xgb_y_pred)

2144.8588299087473
```

```
✓ [85] mean_squared_error(y_test, xgb_y_pred)

8268802.185235631
```

```
✓ [86] r2_score(y_test, xgb_y_pred)

0.6722789165646108
```

```
✓ [87] from math import sqrt
print("RMSE of Linear Regression Model is ",sqrt(mean_squared_error(y_test, xgb_y_pred)))

RMSE of Linear Regression Model is 2875.552500865813
```

The ML algorithm that perform the best was XGBoost Regressor Model with RMSE = 2879